

OPPGAVE 1 (True/False questions) (20 pts, 2 pts for each correct answer, -1 for each wrong answer. Total points will not be less than zero) Mark each of the following sentences either as True or False

1. True or False: Greedy best-first search algorithm is guaranteed to find an optimal path.

ANSWER: False. It takes the first path to goal it finds.

2. True or False: Uniform Cost algorithm is guaranteed to find the optimal solution when all step costs are greater than zero and the branching factor is finite.

ANSWER: True

3. True or False: Let b be the branching factor of a search, d the depth of the solution, and m the maximum depth of the search space. Then the complexity of breadth-first search is b^m .

Answer= false. It is b^d

4. True or false: $(A \Leftrightarrow B) \models (\neg A \vee B)$

ANSWER: True

5. Linear planning is incomplete.

Answer: True, Linear planning is incomplete. Think for example about Sussman's anomaly.

6. Simple "hill climbing" algorithm is perfect to solve constraint satisfaction problems.

ANSWER: False. It can get stuck at local minima and fail to find a solution.

7. A sound logical reasoning process is not necessary in order to pass the Turing test.

ANSWER: True. Humans don't always use sound logical reasoning

8. Depth-first tree search algorithm always expands at least as many nodes as an A* tree search algorithm with admissible heuristic does.

ANSWER: False: a lucky DFS might expand exactly d nodes to reach the goal. A* largely dominates any graph-search algorithm that is guaranteed to find optimal solutions.

9. The set consisting of "mammal" and "non-mammal" categories is both a disjunctive and an exhaustive decomposition of the category "animal".

Answer: True.

10. Explainability is an ethical problem in AI which domain knowledge may help to solve. **Answer: True**

OPPGAVE 2 (First Order Logic) (15pts – 7-3-5)

The Knowledge base has the following sentences for which FOL representations are given below.

Everyone who loves all animals is loved by someone.

Anyone who kills an animal is loved by no one.

Sofie loves all animals.

Either Sofie or CarAccident killed the cat, who is named Kismet.

FOL representations:

1. $\forall x [\forall y [\text{Animal}(y) \Rightarrow \text{Loves}(x,y)]] \Rightarrow [\exists z \text{Loves}(z,x)]$
2. $\forall x [\exists y (\text{Animal}(y) \wedge \text{Kills}(x,y)) \Rightarrow \neg(\exists z \text{Loves}(z,x))]$
3. $\forall x [\text{Animal}(x) \Rightarrow \text{Loves}(\text{Sofie}, x)]$
4. $\text{Kills}(\text{Sofie}, \text{Kismet}) \vee \text{Kills}(\text{CarAccident}, \text{Kismet})$
5. $\text{Cat}(\text{Kismet})$

Query: Did CarAccident kill the cat?

You are going to answer the above query by using resolution by refutation.

1. First convert all FOL sentences into conjunctive normal form (CNF). Show every step in this process.
2. Write down any background knowledge that is needed to solve the problem
3. Apply resolution by refutation and show how the query is answered. Show each and every unification.

ANSWER TO PROBLEM – FOL

CNF:

Everyone who loves all animals is loved by someone.

In FOL

$\forall x \{ [\forall y [\text{Animal}(y) \Rightarrow \text{Loves}(x,y)]] \Rightarrow [\exists z \text{Loves}(z,x)] \}$

Remove Implications

$\forall x \{ [\neg [\forall y \{ \text{Animal}(y) \Rightarrow \text{Loves}(x,y) \}]] \vee [\exists z \text{Loves}(z,x)] \}$

$\forall x \{ [\neg [\forall y \{ \neg \text{Animal}(y) \vee \text{Loves}(x,y) \}]] \vee [\exists z \text{Loves}(z,x)] \}$

Move negation inward

$\forall x \{ [\exists y \{ \neg \neg \text{Animal}(y) \wedge \neg \text{Loves}(x,y) \}] \vee [\exists z \text{Loves}(z,x)] \}$

$\forall x \{ [\exists y \text{Animal}(y) \wedge \neg \text{Loves}(x,y)] \vee [\exists z \text{Loves}(z,x)] \}$

Skolemize

$\forall x \{ [\text{Animal}(F(x)) \wedge \neg \text{Loves}(x,F(x))] \vee [\text{Loves}(G(x),x)] \}$

Drop universal quantifier

$[\text{Animal}(F(x)) \wedge \neg \text{Loves}(x,F(x))] \vee [\text{Loves}(G(x),x)]$

Use distributive law (and get two clauses)

$(\text{Animal}(F(x)) \vee \text{Loves}(G(x),x)) \text{ AND } (\neg \text{Loves}(x,F(x)) \vee \text{Loves}(G(x),x))$

Anyone who kills an animal is loved by no one.

Transfer to FOL

$\forall x \{ [\exists y (\text{Animal}(y) \wedge \text{Kills}(x,y))] \Rightarrow \neg (\exists z \text{Loves}(z,x)) \}$

Remove Implications

$\forall x \{ \neg [\exists y (\text{Animal}(y) \wedge \text{Kills}(x,y))] \vee \neg (\exists z \text{Loves}(z,x)) \}$

Move negations inwards

$\forall x \{ [\forall y \neg \text{Animal}(y) \vee \neg \text{Kills}(x,y)] \vee (\forall z \neg \text{Loves}(z,x)) \}$

Remove quantifiers

$\neg \text{Animal}(y) \vee \neg \text{Kills}(x,y) \vee \neg \text{Loves}(z,x)$

Sofie loves all animals.

FOL form

$\forall x [\text{Animal}(x) \Rightarrow \text{Loves}(\text{Sofie}, x)]$

Remove implications

$\forall x [\neg \text{Animal}(x) \vee \text{Loves}(\text{Sofie}, x)]$

Remove quantifier

$\neg \text{Animal}(x) \vee \text{Loves}(\text{Sofie}, x)$

Either Sofie or CarAccident killed the cat, who is named Kismet.

FOL form

$\text{Kills}(\text{Sofie}, \text{Kismet}) \vee \text{Kills}(\text{CarAccident}, \text{Kismet}), \text{Cat}(\text{Kismet})$

Background knowledge: All cats are animals: $\forall x \text{Cat}(x) \Rightarrow \text{Animal}(x)$

RESOLUTION:

$Cat(Kismet), \neg Cat(x) \vee Animal(x)$

$Unify(Cat(Kismet), \neg Cat(x)) = \{x/Kismet\}$

First line thus resolves to:

$Animal(Kismet)$

$Kills(Sofie, Kismet) \vee Kills(CarAccident, Kismet), \neg Kills(CarAccident, Kismet)$

Resolves to:

$Kills(Sofie, Kismet)$

$\neg Animal(y) \vee \neg Kills(x, y) \vee \neg Loves(z, x), Animal(Kismet)$

$Unify(Animal(Kismet), \neg Animal(y)) = \{y/Kismet\}$

Resolves to:

$\neg Kills(x, Kismet) \vee \neg Loves(z, x),$

$\neg Loves(x, F(x)) \vee Loves(G(x), x), \neg Animal(z) \vee Loves(Sofie, z)$

$Unify(\neg Loves(x, F(x)), Loves(Sofie, z)) = \{x/Sofie, z/F(x)\}$

Resolvent clause is obtained by substituting the unification rule

$Loves(G(Sofie), Sofie) \vee \neg Animal(F(Sofie))$

$Animal(F(x)) \vee Loves(G(x), x), Loves(G(Sofie), Sofie) \vee \neg Animal(F(Sofie))$

$Unify(Animal(F(x)), \neg Animal(F(Sofie))) = \{x/Sofie\}$

Resolvent clause is obtained by substituting the unification rule

$Loves(G(Sofie), Sofie)$

$\neg Kills(x, Kismet) \vee \neg Loves(z, x), Loves(G(Sofie), Sofie)$

$Unify(\neg Loves(z, x), Loves(G(Sofie), Sofie)) = \{x/Sofie, z/G(Sofie)\}$

Resolvent clause is obtained by substituting the unification rule

$\neg Loves(G(Sofie), Sofie)$

$\neg Loves(G(Sofie), Sofie), Loves(G(Sofie), Sofie)$

Resolvent clause is empty. Proof succeeded

OPPGAVE 3 - A* search algorithm (13 pts, 1-4-4-4)

1. Apply A* algorithm for graphs on the graph in the following figure (Figure 1). Write down the nodes in the order they are expanded. A is the start node, and G is the goal node.
2. What is the returned path? Is it optimal? explain your answer on an example from the given graph in the figure.
3. Modify the pseudocode in the figure (Figure 2) for A* algorithm for graph search so that it guarantees to find the optimal solution with heuristic values given in Figure 1. Write down the pseudocode in a separate paper starting from the sentence just before your modification starts, and ending with the sentence right after your last change. That is, you don't need to write the whole pseudocode, only the part you modified, plus a single/one original sentence before and one after your modified sentences.
4. Assume that a search tree has heuristic values which enable the A* algorithm presented in the pseudocode (the unmodified version in figure 2) to find an optimal path to the goal. Would the A* algorithm still be guaranteed to find the minimal path to the goal if there are negative transition costs? This question is general, not about the problem presented in Figure 1.




Figure 1 Graph for the question. A is the start node, and G is the goal node.


1. `Start.g = 0;`
2. `Start.h = heuristic(Start)`
3. `FRONTIER = {Start}`
4. `CLOSED = {empty set}`
5. `WHILE FRONTIER is not empty`
 6. `N = FRONTIER.popLowestF()`
 7. `IF state of N= GOAL RETURN N`
 8. `add N to CLOSED`
 9. `FOR all children M of N not in CLOSED:`
 10. `M.parent = N`
 11. `M.g = N.g + cost(N,M)`
 12. `M.h = heuristic(M)`
 13. `add M to FRONTIER`
 14. `ENDFOR`
15. `ENDWHILE`

Figure 2 .Pseudocode for A* algorithm.

ANSWER to PROBLEM (A* alg.)


1. Expanded nodes: ACBDG.

2. Returned path: ACG. Not optimal. Because of B is not consistent. $H(B) > h(C) + c(B,C)$ C is expanded with path cost 6, but then it appears as the child of B again with a path cost 5 ($3+2$) this time but since it is in "closed" is not re-expanded.
3. The pseudocode is changed so that a node in "closed" can be re-expanded (by putting it back into "frontier". This can be done by removing the "not in CLOSED" part of Sentence 9. Alternatively, the path cost ($M.g+M.h$) of the new encounter of the node (M) is compared with path cost of M in the "closed" and it is put back to frontier only if its new path cost is smaller.
4. No. See the example below (cost 1 versus -99)



OPPGAVE 4 - CONSTRAINT SATISFACTION (12 pts, 2-4-2-4)

Consider the following constraint graph (in the figure) for a graph coloring problem where the constraints mean that the connected nodes cannot have the same color. The variables are shown inside the nodes while the domains are shown next to each variable node.



1. What are the domains after a full constraint propagation using an arc consistency algorithm?
2. Show the sequence of variable assignments during a pure backtracking search (don't assume that propagation above has been done). Assume that the variables are examined in numerical order and the values are assigned in the order shown next to each node. Show assignments by writing the variable number and the letter for the value, e.g. 5R, 2G.
3. Describe how forward checking works.
4. This time you'll apply backtracking search with forward checking. Use the same ordering convention for variables and values as above. Show the sequence of variable assignments during backward search with forward checking. Again,

show assignments by writing the number of variables followed by the letter for the value.

ANSWER to PROBLEM (Constaint Satisfaction)

1. 1={R}

2={G,B}

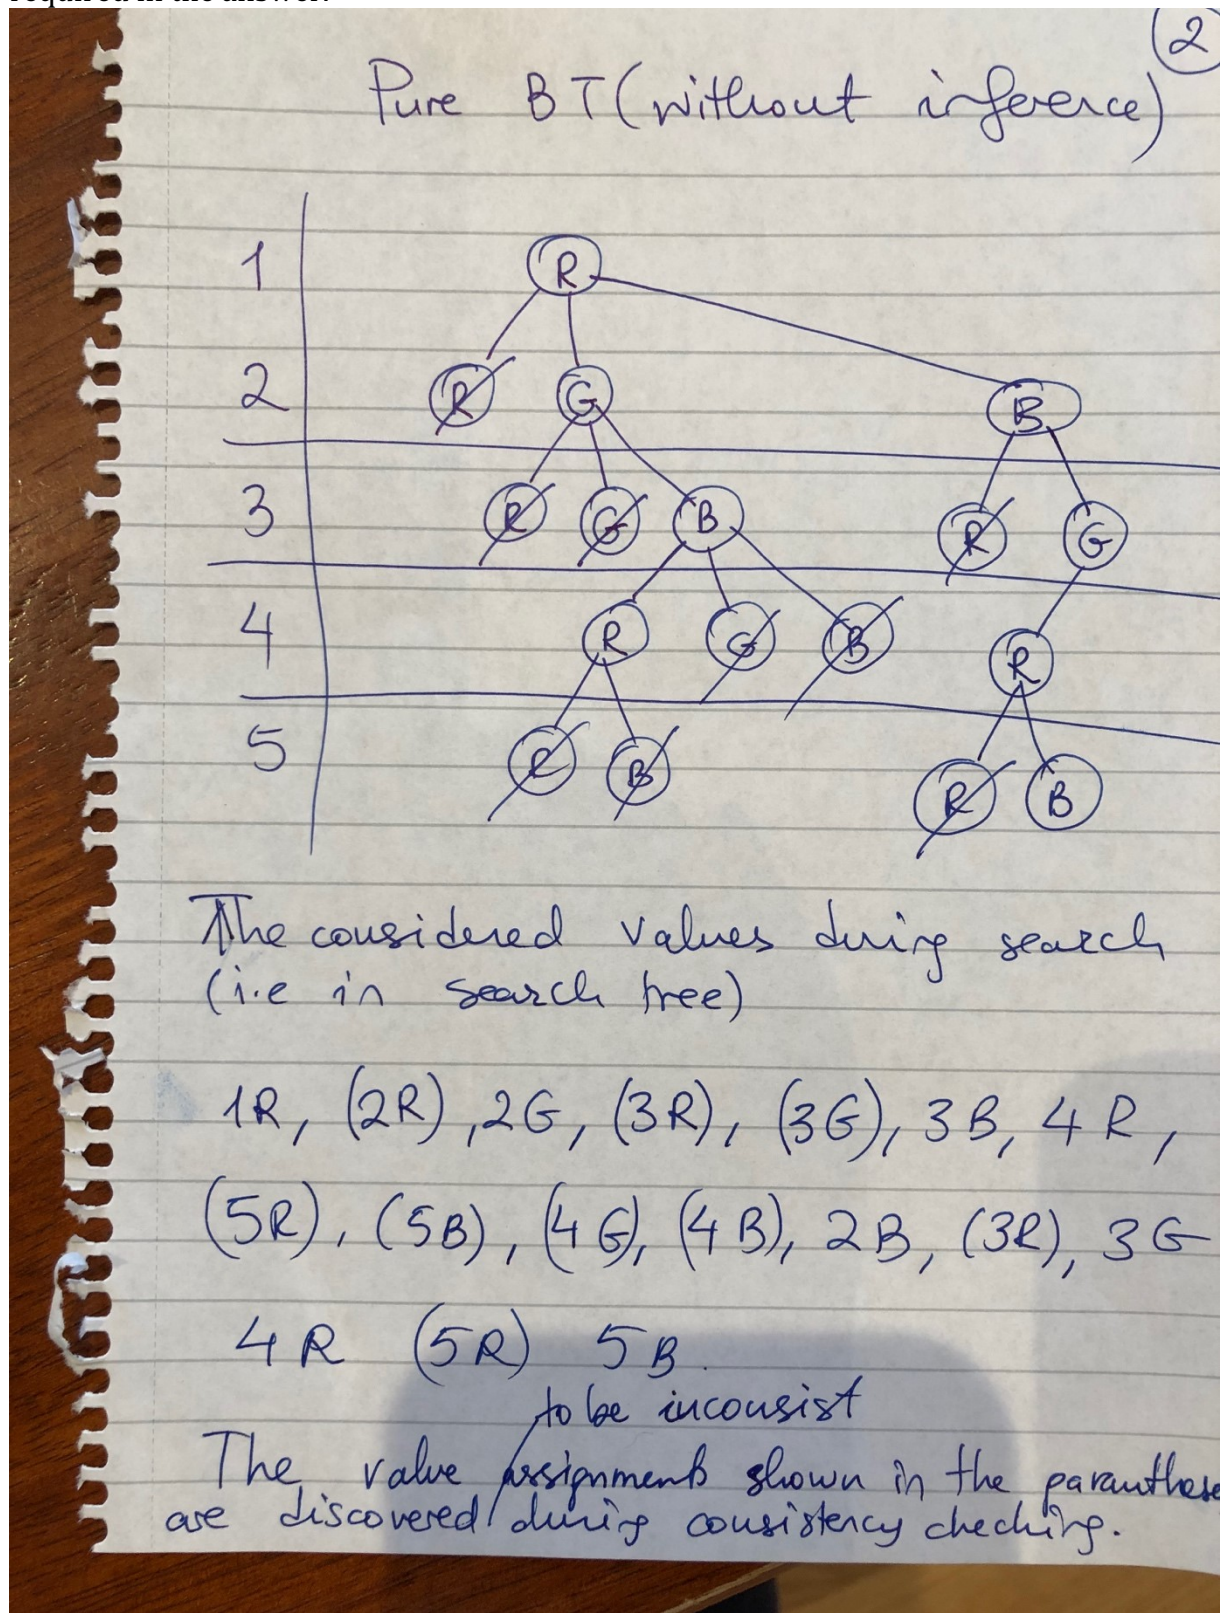
3={G,B}

4={R,G,B}

5={R,B}

2. 1R, 2R, 2G, 3R,3G, 3B, 4R, 5R, 5B, 4G, 4B, 2B, 3R, 3G, 4R, 5R, 5B. The tried but not assigned values during backtracking search are shown in parentheses: 1R, (2R), 2G, (3R), (3G), 3B, 4R, (5R), (5B), (4G), (4B), 2B, (3R), 3G, 4R, (5R), 5B. The figure below shows all value considerations. Full 4 points have been given also to answers that don't include the inconsistent values (i.e., the ones in the parenthesis), i.e., 1R, 2G, 3B, 4R, 2B, 3G, 4R, 5B is also accepted as correct answer. So both 1R, 2R, 2G, 3R,3G, 3B, 4R, 5R, 5B, 4G, 4B, 2B, 3R, 3G, 4R, 5R, 5B . and 1R, 2G, 3B, 4R, 2B, 3G, 4R, 5B are accepted as correct answer.

The following and the next figure are in order to show/explain you, they are not required in the answer.



3. Forward checking does check only 1 step (immediate neighbours not assigned yet) after assignment of a variable.

4. 1R, 2G, 3B, 4R, 2B, 3G, 4R, 5B. No point is given if the answer is the result path only (i.e., 1R, 2B, 3G, 4R, 5B)

①

BT with FC.

	1	2	3	4	5
initial	R	RGB	RGB	RGB	RB
1 R	<u>R</u>	<u>GB</u>	<u>GB</u>	<u>RGB</u>	<u>RB</u>
2 G	<u>R</u>	<u>G</u>	<u>B</u>	<u>RB</u>	<u>RB</u>
3 B	<u>R</u>	<u>G</u>	<u>B</u>	<u>R</u>	<u>R</u>
4 R	<u>R</u>	<u>G</u>	<u>B</u>	<u>R</u>	X Backed
2 B	<u>R</u>	<u>B</u>	<u>G</u>	<u>RG</u>	<u>RB</u>
3 G	<u>R</u>	<u>B</u>	<u>G</u>	<u>R</u>	<u>B</u>
4 R	<u>R</u>	<u>B</u>	<u>G</u>	<u>R</u>	<u>B</u>
5 B					OK!

So, the value assignments:

1 ~~R~~, 2G, 3B, 4R (backtrack), 2B
 3G, ~~4R~~, 5B.

OPPGAVE 5 – Game Theory (10 pts, equal points for each question))

Consider the game for which the payoff matrix is shown in the following figure.

		agent2	
		G	NG
agent1	H	8, 0	3, 1
	NH	4, 4	2, 3

Figure 3 Payoff Matrix.

1. Identify any *dominated* strategy. Explain your answer.
2. Find the *Nash equilibrium*. What are the equilibrium payoffs, i.e., values for each agent?
3. Are there any *pareto optimal* joint actions? If any exists, what are they?
4. Explain (in general, not for this particular problem) why a *social welfare maximizing* joint action profile is also pareto optimal.

Answer:

1. First notice, neither of Agent2's strategies are dominated since,

$u_2(H, G) = 0 < 1 = u_2(H, NG)$ and $u_2(NH, G) = 4 > 3 = u_2(NH, NG)$
H strictly dominates NH for Agent1 since,
 $u_1(H, G) = 8 > 4 = u_1(NH, G)$
and $u_1(H, NG) = 3 > 2 = u_1(NH, NG)$.

2. NE is (H, NG) which yields a payoff of (3, 1)

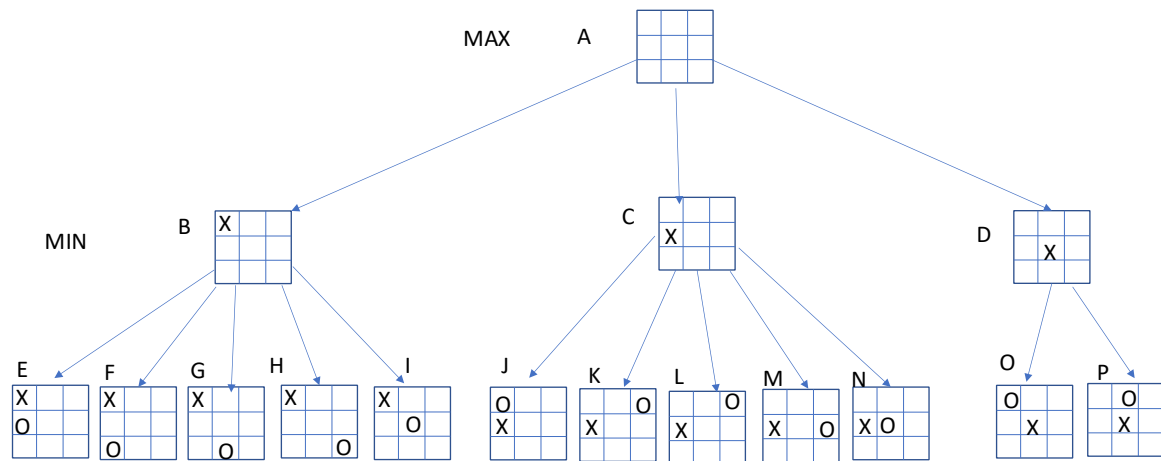
3. (H,G) and (NH, G) are pareto optimal

4. Social welfare maximizing profile means that the sum of utilities of all agents are highest for this action profile. This means that it is not possible to increase utilities of both agents at the same time.

OPPGAVE 6 - ADVERSARIAL SEARCH (10 pts, 3-3-4)

Consider a tic-tac-toe game on a 3x3 grid (see the figure below) where players MAX and

MIN take turns marking the spaces in a 3×3 grid by placing their X's and O's, respectively. The player who succeeds in placing three of their marks in a horizontal, vertical, or diagonal row is the winner.



The values of the terminal nodes have not been provided. You will compute these values using an evaluation function e . Function e uses a heuristic that estimates the value of each terminal node according to the following formula:

$e(\text{node}) = E1 - E2$ where

$E1$ = sum of the number of rows, columns and diagonals that are possible winning situations for Max and,

$E2$ = sum of the number of rows, columns and diagonals that are possible winning situations for Min.

The following figure shows examples for computing the values of some hypothetical nodes:





Figure : Example on estimation the values of terminal nodes

1. Compute and write down the values of each terminal node node, i.e of E,, P.

- Using the *Minimax* algorithm find and write down the values of the remaining nodes(A,...D) in the search tree. Which action will Max play in this game, the action whose outcome is B, C or D?
- Is it possible to prune any nodes using the *Alpha-Beta pruning* algorithm? If there are any, write down the prunable node(s) - i.e., write the letter for the node.

ANSWER to ADVERSARIAL SEARCH

- See figure for MINIMAX alg. Best action is the one with outcome D.
- K,L, M,N can be pruned. Answers excluding K or excluding L are also correct since the two are equal.



OPPGAVE 7 - SHORT QUESTIONS (20 pts, 2 pts for each question. Plus 2 points for correct answer to the “bonus”/voluntary question.)

1. How is the goal information represented in simple reflex agents?

Answer: the goal is implicit in the condition-action rule (goals are “designed in”)

2. Is the PL sentence $((P \rightarrow Q) \wedge Q) \rightarrow P$ valid, unsatisfiable or satisfiable? Justify your answer.

ANSWER: Satisfiable. When $P=T$ and $Q=T$ the sentence is true, but when $P=F$ and $Q=T$ the sentence is false.

3. Translate the following sentence into predicate logic:
“Any person who has an umbrella is not wet”

ANSWER: $\forall x [(IsPerson(x) \Rightarrow \exists y (Has(x, y) \wedge IsUmbrella(y))) \Rightarrow \neg IsWet(x)]$.

4. Translate the following sentence into predicate logic:
“John has at least two daughters.”

ANSWER: $\exists x, y Daughter(x, John) \wedge Daughter(y, John) \wedge \neg (x=y)$

5. Assume a Hill Climbing algorithm that aims to find the best state according to a heuristic cost function. Does it try to find the global minimum or the global maximum?

Answer: Global minimum

6. It has been suggested that the first phase of *GraphPlan* be used as a heuristic function for forward search in the following way: Given a state s and goal g , run the graph-construction phases of *GraphPlan* until all the components are represented and not mutex in the last layer. Let n be the number of action layers in the graph. We will let n be the heuristic value for s . Is this an admissible heuristic? Explain your answer.

ANSWER: Yes, this is an admissible heuristic because it will always

underestimate the distance to a solution (no solution can be nearer than the first layer where all of the solution propositions are no mutex)

7. If *GraphPlan* terminates with a successful, 3-action plan in the first iteration, what constraints are there on the order in which the actions must be executed?

ANSWER: There are no constraints on what order the actions must be executed. They are all in the same layer, which indicates that they can be performed in parallel.

8. Does *Regression Planning* work in a forward or backward manner?

ANSWER: Backward

9. You will represent the concept of “student” using a *frame-based* knowledge representation language. You want the age of a student to be computed on the basis of her birth year and the current year. How would you represent this in a slot of “student” frame?

Answer: In the “age” slot, and as a demon, a procedure/function that computes the age.

Name: student
isa: person
age: (calculate-age)
... ..
...

10. Assume a version of the original vacuum cleaner agent in the textbook. 10% of the time the SUCK action of this one does not clean the floor if it is dirty and even may deposit dirt on the dirt on the floor if the floor is clean. Classify this environment with respect to each of the following dimensions:

Sequential/Episodic, deterministic/stochastic, and dynamic versus static.

Answer: sequential, stochastic, static

11. BONUS QUESTION: Assume that *Simulated Annealing* search algorithm starts from a state S_0 in the middle of a large plateau. That is, the values of all states on the plateau are exactly the same. Assume also that in the first step the random neighbor we picked is S_1 , which has the same value as S_0 . Will Simulated annealing move to S_1 ? Explain your answer mathematically (i.e. using a formula). No points will be given otherwise.

ANSWER: True. Since $P(S_0 \rightarrow S_1) = \exp(\Delta E/T)$ and

$\Delta E = |\text{VALUE}(S_0) - \text{VALUE}(S_1)|$, we see that with probab. 1 we will move to S_1 .

